

# Chapter 8: Satisfiability and Model Construction

Source A: Foundations of Artificial Intelligence by J. Boedecker, W. Burgard, F. Hutter, B. Nebel (Uni. Freiburg)

# Motivation

SAT solving is the best available technology for practical solutions to many NP-hard problems Formal verification

- Verification of software
  - Ruling out unintended states (null-pointer exceptions, etc.)
  - Proving that the program computes the right solution
- Verification of hardware (Pentium bug, etc)

Practical approach: encode into SAT & exploit the rapid progress in SAT solving

- Solving CSP instances in practice
- Solving graph coloring problems in practice

# Contents

- The SAT Problem (Erfuellbarkeitsproblem der Aussagenlogik)
- Davis-Putnam-Logemann-Loveland (DPLL) Procedure
- "Average" Complexity of the Satisfiability Problem
- Local Search Procedures
- State of the Art

# Logical deduction vs. satisfiability

## Propositional Logic - typical algorithmic questions:

- Logical deduction
  - Given: A logical theory (set of propositions)
  - Question: Does a proposition logically follow from this theory?
  - Reduction to unsatisfiability, which is coNP-complete (complementary to NP problems)
- Satisfiability of a formula (SAT)
  - Given: A logical theory
  - Wanted: Model of the theory
  - Example: Configurations that fulfill the constraints given in the theory
  - Can be "easier" because it is enough to find one model
  - Possible Application: Find software bugs, i.e. find a state that is not desired

# The Satisfiability Problem (SAT)

- Given
  - Propositional formula  $\phi$  in CNF
- Wanted
  - Model of  $\phi$ .
  - or proof, that no such model exists.

SAT can be formulated as a Constraint-Satisfaction-Problem ( $\rightarrow$  search):

- CSP-Variables = Symbols of the alphabet
- Domain of values =  $\{T, F\}$
- Constraints given by clauses

# The DPLL algorithm

The DPLL algorithm (Davis, Putnam, Logemann, Loveland, 1962) corresponds to backtracking with inference in CSPs:

- Recursive call DPLL ( $\Delta$ ,  $l$ ) with
  - $\Delta$ : set of clauses
  - $l$ : variable assignment
- Result: satisfying assignment that extends  $l$  or "unsatisfiable" if no such assignment exists.
  - First call by DPLL( $\Delta$ ,  $\emptyset$ )
- Inference in DPLL:
  - Simplify: if variable  $v$  is assigned a value  $d$ , then all clauses containing  $v$  are simplified immediately (corresponds to forward checking)
  - Variables in unit clauses (= clauses with only one variable) are immediately assigned (corresponds to minimum remaining values ordering in CSPs)

# The DPLL Procedure

## DPLL Function

Given a set of clauses  $\Delta$  defined over a set of variables  $\Sigma$ , return "satisfiable" if  $\Delta$  is satisfiable. Otherwise return "unsatisfiable".

- 1. If  $\Delta = \emptyset$  return "satisfiable"
- 2. If  $\epsilon \in \Delta$  return "unsatisfiable"
- 3. Unit-propagation Rule: If  $\Delta$  contains a unit-clause  $C$ , assign a truth-value to the variable in  $C$  that satisfies  $C$ , simplify  $\Delta$  to  $\Delta'$  and return  $\text{DPLL}(\Delta')$ .
- 4. Splitting Rule: Select from  $\Sigma$  a variable  $v$  which has not been assigned a truth-value. Assign one truth value  $t$  to it, simplify  $\Delta$  to  $\Delta'$  and call  $\text{DPLL}(\Delta')$ 
  - a. If the call returns "satisfiable", then return "satisfiable".
  - b. Otherwise assign the other truth-value to  $v$  in  $\Delta$ , simplify to  $\Delta''$  and return  $\text{DPLL}(\Delta'')$ .

## Example (1)

$$\Delta = \{\{a, b, \neg c\}, \{\neg a, \neg b\}, \{c\}, \{a, \neg b\}\}$$

1. Unit-propagation rule:  $c \mapsto T$   
 $\{\{a, b\}, \{\neg a, \neg b\}, \{a, \neg b\}\}$
2. Splitting rule:

2a.  $a \mapsto F$   
 $\{\{b\}, \{\neg b\}\}$

3a. Unit-propagation rule:  
 $b \mapsto T$   
 $\{\square\}$



## Example (1)

$$\Delta = \{\{a, b, \neg c\}, \{\neg a, \neg b\}, \{c\}, \{a, \neg b\}\}$$

1. Unit-propagation rule:  $c \mapsto T$   
 $\{\{a, b\}, \{\neg a, \neg b\}, \{a, \neg b\}\}$

2. Splitting rule:

2a.  $a \mapsto F$   
 $\{\{b\}, \{\neg b\}\}$

2b.  $a \mapsto T$   
 $\{\{\neg b\}\}$

3a. Unit-propagation rule:  
 $b \mapsto T$   
 $\{\square\}$

3b. Unit-propagation rule:  $b \mapsto F$   
 $\{\}$

## Example (2)

$$\Delta = \{\{a, \neg b, \neg c, \neg d\}, \{b, \neg d\}, \{c, \neg d\}, \{d\}\}$$

- 1. Unit-propagation rule:  $d \rightarrow T$ 
  - $\{\{a, \neg b, \neg c\}, \{b\}, \{c\}\}$
- 2. Unit-propagation rule:  $b \rightarrow T$ 
  - $\{\{a, \neg c\}, \{c\}\}$
- 3. Unit-propagation rule:  $c \rightarrow T$ 
  - $\{\{a\}\}$
- 4. Unit-propagation rule:  $a \rightarrow T$ 
  - $\{\}$

# Properties of DPLL

- DPLL is complete, correct, and guaranteed to terminate.
- DPLL constructs a model, if one exists.
- In general, DPLL requires exponential time (splitting rule!)
  - → Heuristics are needed to determine which variable should be instantiated next and which value should be used.
- DPLL is polynomial on Horn clauses (see next slides).
- In current SAT competitions, DPLL-based procedures have shown the best performance.

## DPLL on Horn Clauses (0)

- Horn Clauses constitute an important special case, since they require only polynomial runtime of DPLL.
- Definition: A Horn clause is a clause with maximally one positive literal E.g.,  $\neg A_1 \vee \dots \vee \neg A_n \vee B$  or  $\neg A_1 \vee \dots \vee \neg A_n$  ( $n = 0$  is permitted).
- Equivalent representation:  $\neg A_1 \vee \dots \vee \neg A_n \vee B \Leftrightarrow \bigwedge_i A_i \Rightarrow B$ 
  - $\rightarrow$  Basis of logic programming (e.g., PROLOG)

# DPLL on Horn Clauses (1)

## Note

- 1. The simplifications in DPLL on Horn clauses always generate Horn clauses
- 2. If the first sequence of applications of the unit propagation rule in DPLL does not lead to termination, a set of Horn clauses without unit clauses is generated
- 3. A set of Horn clauses without unit clauses and without the empty clause is satisfiable, since
  - All clauses have at least one negative literal (since all non-unit clauses have at least two literals, where at most one can be positive (Def. Horn))
  - Assigning *false* to all variables satisfies formula

## DPLL on Horn Clauses (2)

... still note

- 4. It follows from 3.:
  - a. every time the splitting rule is applied, the current formula is satisfiable
  - b. every time, when the wrong decision (= assignment in the splitting rule) is made, this will be immediately detected (e.g., only through unit propagation steps and the derivation of the empty clause).
- 5. Therefore, the search trees for  $n$  variables can only contain a maximum of  $n$  nodes, in which the splitting rule is applied (and the tree branches).
- 6. Therefore, the size of the search tree is only polynomial in  $n$  and therefore the running time is also polynomial.

## "Average" Complexity of the Satisfiability Problem

"Average" Complexity of the Satisfiability Problem

# How Good is DPLL in the Average Case?

- We know that SAT is NP-complete, i.e., in the worst case, it takes exponential time.
- This is clearly also true for the DPLL-procedure. → Couldn't we do better in the average case?
- For CNF-formulae, in which the probability for a positive appearance, negative appearance and non-appearance in a clause is  $1/3$ , DPLL needs on average quadratic time (Goldberg 79)!
  - → The probability that these formulae are satisfiable is, however, very high.




## Phase Transitions ...

- Conversely, we can, of course, try to identify hard to solve problem instances.
- Cheeseman et al. came up with the following plausible conjecture:
  - All NP-complete problems have at least one order parameter and the hard to solve problems are around a critical value of this order parameter. This critical value (a *phasetransition*) separates one region from another, such as over-constrained and under-constrained regions of the problem space.
  - over-constrained: you have more and more clauses and hence it becomes more and more difficult to find a solution
  - under-constrained: you have much more variables than constraints and hence a solution is easy to find
- Confirmation for graph coloring and Hamiltonian path ..., later also for other NP-complete problems.


# Phase Transitions with 3-SAT

- Constant clause length model (Mitchell et al.): Clause length  $k$  is given. Choose variables for every clause  $k$  and use the complement with probability 0.5 for each variable.
- Phase transition for 3-SAT with a clause/variable ratio of approx. 4.3:



# Empirical Difficulty

- The Davis-Putnam (DPLL) Procedure shows extreme runtime peaks at the phase transition:
- Note: Hard instances can exist even in the regions of the more easily satisfiable/unsatisfiable instances!



# Notes on the Phase Transition

- When the probability of a solution is close to 1 (under-constrained), there are many solutions, and the first search path of a backtracking search is usually successful.
- If the probability of a solution is close to 0 (over-constrained), this fact can usually be determined early in the search.
- In the phase transition stage, there are many near successes ("close, but no cigar")
  - → (limited) possibility of predicting the difficulty of finding a solution based on the parameters
  - → (search intensive) benchmark problems are located in the phase region (but they have a special structure)

# Local Search Procedures

## Local Search Procedures

# Local Search Methods for Solving Logical Problems

- In many cases, we are interested in finding a satisfying assignment of variables (example CSP), and we can sacrifice completeness if we can "solve" much larger instances this way.
- Standard process for optimization problems: Local Search
  - Based on a (random) configuration
  - Through local modifications, we hope to produce better configurations
  - → Main problem: local maxima

# Dealing with Local Maxima

- As a measure of the value of a configuration in a logical problem, we could use the number of satisfied constraints/clauses.
- At first glance, local search seems inappropriate, considering that we want to find a global maximum (all constraints/clauses satisfied).
- However:
  - By restarting and/or injecting noise, we can often escape local maxima.
  - Local search can perform very well for SAT solving

# A pioneering local search method for SAT: GSAT (1993)


## Procedure GSAT

```
INPUT: a set of clauses  $\alpha$ ,  $Max - Flips$ , and  $Max - Tries$ 
OUTPUT: a satisfying truth assignment of  $\alpha$ , if found
begin
  for i := 1 to  $Max - Tries$ 
    T := a randomly-generated truth assignment
    for j := 1 to  $Max - Flips$ 
      if T satisfies  $\alpha$  then return T
      v := a propositional variable such that a change in its truth ...
      ...assignment gives the largest increase in
the number of clauses of  $\alpha$  that are satisfied by T
      T := T with the truth assignment of v reversed
    end for
  end for
  return "no satisfying assignment found"
end
```



# The Search Behavior of GSAT

- In contrast to many other local search methods, we must also allow sideways movements!
- Most time is spent searching on plateaus.



# State of the Art

## State of the Art

# Practical Improvements of DPLL Algorithms

## Clause Learning

- Consider an exemplary SAT problem
  - 26 variables  $A, \dots, Z$
  - Amongst many other clauses, we have  
 $\{(\neg A, Y, Z)\}, \{(\neg A, \neg Y, Z)\}, \{(\neg A, Y, \neg Z)\}, \{(\neg A, \neg Y, \neg Z)\}$
  - We'll branch on variables in lexicographic order and try true first
- What will happen?
  - There is no satisfying assignment to the clauses above when  $A = T$
  - For each assignment to variables  $B, \dots, X$ , we'll have to rediscover this fact
  - Rather: reason about the variables that led to a conflict:  $A, Y$  and  $Z$
  - We can 'Learn' (here: logically infer) a new clause:  $\neg A$
  - Leads to conflict-directed clause learning (CDCL)
- Intelligent Backjumping
  - Closely related to clause learning
  - Jump back to the branching decision responsible for a conflict

# Practical Improvements of SAT Algorithms

Both for DPLL/CDCL algorithms and local search algorithms

- Efficient data structures, indexing, etc
- Engineering ingenious heuristics
- Meta-algorithmic advances
  - Automated parameter tuning and algorithm configuration
  - Selection of the best-fitting algorithm based on instance characteristics
  - Selection of the best-fitting parameters based on instance characteristics
  - Use of machine learning to pinpoint what factors most affects performance

# The Current State of the Art

- SAT competitions since beginning of the 90s
- Current SAT competitions (<http://www.satcompetition.org/>):
  - Largest "industrial" instances:  $> 10,000,000$  variables
- Complete solvers dominate handcrafted and industrial tracks
- Incomplete local search solvers best on random satisfiable instances
- Best solvers use meta-algorithmic methods, such as algorithm configuration, selection, etc.

# Concluding Remarks

- SAT solving very prominently uses resolution
- DPLL: combines resolution and backtracking
  - Very efficient implementation techniques
  - Good branching heuristics
  - Clause learning
- Incomplete randomized SAT-solvers
  - Perform best on random satisfiable problem instances
- State of the art
  - Typically obtained by automatic algorithm configuration & selection