

Prof. Dr. Matthias Schott


- 1 What is Action Planning?
- 2 Planning Formalisms
- 3 Basic Planning Algorithms
- 4 Computational Complexity

Planning

- Planning is the process of generating (possibly partial) representations of future behavior prior to the use of such plans to constrain or control that behavior.
- The outcome is usually a set of actions, with temporal and other constraints on them, for execution by some agent or agents.


Planning Tasks

- Given a current state, a set of possible actions, a specification of the goal conditions, which plan transforms the current state into a goal state?



[illegible]

- Given a road map, and a number of trucks and airplanes, make a plan to transport objects from their start to their goal destinations.



Action Planning is not ...

- Problem solving by search, where we describe a problem by a state space and then implement a program to search through this space
 - in action planning, we specify the problem declaratively (using logic) and then solve it by a general planning algorithm
- Program synthesis, where we generate programs from specifications or examples
 - in action planning we want to solve just one instance and we have only very simple action composition (i.e., sequencing, perhaps conditional and iteration)
- Scheduling, where all jobs are known in advance and we only have to fix time intervals and machines
 - instead we have to find the right actions and to sequence them
- Of course, there is interaction with these areas!

Domain-Independent Action Planning

- Start with a declarative specification of the planning problem
- Use a domain-independent planning system to solve the planning problem
- Domain-independent planners are generic problem solvers
- Issues:
 - Good for evolving systems and those where performance is not critical
 - Running time should be comparable to specialized solvers
 - Solution quality should be acceptable
 - →... at least for all the problems we care about

- | | | | |
|---------------------------|----------------------------|---------------|--------|
| Prof. Dr. Matthias Schott | Introduction to AI (Ref A) | June 20, 2021 | 9 / 66 |
|---------------------------|----------------------------|---------------|--------|

Operators, Actions And State Change

- Operator: $a = \langle para, pre, eff \rangle$, i.e. parameter, precondition and effect
 - with $para \subseteq V, pre \subseteq \Sigma_{S,V}, eff \subseteq \Sigma_{S,V} \cup \neg\Sigma_{S,V}$ (element-wise negation) and all variables in pre and eff are listed in $para$.
- Also: $pre(o), eff(o)$.
 - $eff^+ =$ positive effect literals
 - $eff^- =$ negative effect literals
- Operator instance or action: Operator with empty parameter list (instantiated schema!)
- State change induced by action:

$$App(S, o) = \begin{cases} S \cup eff^+(o) - \neg eff^-(o) & \text{if } pre(o) \subseteq S \text{ \& } eff(o) \text{ is cons.} \\ \text{undefined} & \text{otherwise} \end{cases}$$

Example Formalization: Logistics

- Logical atoms: $at(O, L)$, $in(O, V)$, $airconn(L1, L2)$, $street(L1, L2)$, $plane(V)$, $truck(V)$ (i.e. essentially we define predicates here)
 - capital letters here are variables
- Load into truck (define an operation): load
 - Parameter list: (O, V, L)
 - Precondition: $at(O, L), at(V, L), truck(V)$
 - Effects: $\neg at(O, L), in(O, V)$
- Drive operation: drive
 - Parameter list: $(V, L1, L2)$
 - Precondition: $at(V, L1), truck(V), street(L1, L2)$
 - Effects: $\neg at(V, L1), at(V, L2)$
- Some constant symbols: $v1, s, t$ with $truck(v1)$ and $street(s, t)$
- Action: $drive(v1, s, t)$

- A plan \triangle is a sequence of actions
- State (S) resulting (Res) from executing a plan (\triangle):

$$\begin{aligned} Res(S, \langle \rangle) &= S \\ Res(S, (o; \Delta)) &= \begin{cases} Res(App(S, o), \Delta) & \text{if } App(S, o) \\ & \text{is defined} \\ \text{undefined} & \text{otherwise} \end{cases} \end{aligned}$$

- Plan Δ is successful or solves a planning task if $Res(I, \Delta)$ is defined and $G \subseteq Res(I, \Delta)$.

A Small Logistics Example

Initial state: $S = \left\{ \begin{array}{l} at(p1, c), at(p2, s), at(t1, c), \\ at(t2, c), street(c, s), street(s, c) \end{array} \right\}$

Goal: $G = \{ at(p1, s), at(p2, c) \}$

Successful plan: $\Delta = \langle load(p1, t1, c), drive(t1, c, s), \\ unload(p1, t1, s), load(p2, t1, s), \\ drive(t1, s, c), unload(p2, t1, c) \rangle$

- Other successful plans are, of course, possible

Beyond STRIPS

- Even when keeping all the restrictions of classical planning, one can think of a number of extensions of the planning language.
- General logical formulas as preconditions
- Conditional effects: Effects that happen only if some additional conditions are true. For example, when pressing the accelerator pedal, the effects depends on which gear has been selected (no, reverse, forward).
- Multi-valued state variables: Instead of 2-valued Boolean variables, multi-valued variables could be used
- Numerical resources: Resources (such as fuel or time) can be effected and be used in preconditions
- Durative actions: Actions can have duration and can be executed concurrently
- Axioms/Constraints: The domain is not only described by operators, but also by additional laws

Downloaded from <http://ajph.org/> on November 10, 2015

PDDL Logistics Example

```
(define (domain logistics)
  (:types truck airplane - vehicle
           package vehicle - physobj
           airport location - place
           city place physobj - object)
  ...
  (:predicates (in-city ?loc - place ?city - city)
                (at ?obj - physobj ?loc - place)
                (in ?pkg - package ?veh - vehicle))
  ...
  (:action LOAD-TRUCK
    :parameters (?pkg - package ?truck - truck ?loc - place)
    :precondition (and (at ?truck ?loc) (at ?pkg ?loc))
    :effect (and (not (at ?pkg ?loc)) (in ?pkg ?truck))) ...)
```

- variable notation: *?var*

```
(define (state)
  ...)
```

Content


- 3 Basic Planning Algorithms

Basic Planning Algorithms


- Planning Problems as Transition Systems
- We can view planning problems as searching for goal nodes in a large labeled graph (transition system)
- Nodes are defined by the value assignment to the fluents = states
- Labeled edges are defined by actions that change the appropriate fluents
- Use graph search techniques to find a (shortest) path in this graph!
- Note: The graph can become huge: 50 Boolean variables lead to $2^{50} = 10^{15}$ states
- → Create the transition system on the fly and visit only the parts that are necessary

Transition System: Searching Through the State Space


(1/4)




(2/4)



(3/4)



(4/4)



Search through transition system starting at initial state

- Need to use some search strategy. Progression planning can be easily extended to more expressive planning languages


Progression Planning: Example (1/3)

$$\mathcal{S} = \{a, b, c, d\},$$

$$\mathbf{O} = \begin{cases} o_1 = \langle \emptyset, \{a, b\}, \{-b, c\} \rangle, \\ o_2 = \langle \emptyset, \{a, b\}, \{-a, -b, d\} \rangle, \\ o_3 = \langle \emptyset, \{c\}, \{b, d\} \rangle, \end{cases}$$

$$\mathbf{I} = \{a, b\}$$

$$\mathbf{G} = \{b, d\}$$




Progression Planning: Example (2/3)

$$\mathcal{S} = \{a, b, c, d\},$$

$$\mathbf{O} = \{ \begin{array}{l} o_1 = \langle \emptyset, \{a, b\}, \{\neg b, c\} \rangle, \\ o_2 = \langle \emptyset, \{a, b\}, \{\neg a, \neg b, d\} \rangle, \\ o_3 = \langle \emptyset, \{c\}, \{b, d\} \rangle, \end{array}$$


$$\mathbf{I} = \{a, b\}$$

$$\mathbf{G} = \{b, d\}$$



$$\mathbf{O} = \begin{cases} o_1 = \langle \emptyset, \{a, b\}, \{-b, c\} \rangle, \\ o_2 = \langle \emptyset, \{a, b\}, \{-a, -b, d\} \rangle, \\ o_3 = \langle \emptyset, \{c\}, \{b, d\} \rangle, \end{cases}$$

$$\mathbf{G} = \{b, d\}$$



Regression Planning: Backward Search

Search through transition system starting at goal states. Consider sets of states, which are described by the atoms that are necessarily true in them

- 1) Initialize partial plan $\Delta := \langle \rangle$ and set $S := G$
- 2) Test whether we have reached the unique initial state already: $I \supseteq S$? If so, return plan Δ .
- 3) Select one action o_i which does not make (sub-)goals false ($S \cap \neg eff^-(o_i) = 0$) and
 - compute the regression of the description S through o_i :

$$S := S - eff^+(o_i) \cup pre(o_i)$$
 - extend plan $\Delta := \langle o_i, \Delta \rangle$, and continue with step 2.

Need to use some search strategy! Regression becomes much more complicated, if e.g. conditional effects are allowed. Then the result of a regression can be a general Boolean formula

Regression Planning: Example (1/4)

$$\mathcal{S} = \{a, b, c, d, e\},$$

$$\mathbf{O} = \begin{cases} o_1 = \langle \emptyset, \{b\}, \{\neg b, c\} \rangle, \\ o_2 = \langle \emptyset, \{e\}, \{b\} \rangle, \\ o_3 = \langle \emptyset, \{c\}, \{b, d, \neg e\} \rangle, \end{cases}$$

$$\mathbf{I} = \{a, b\}$$

$$\mathbf{G} = \{b, d\}$$

$\{b, d\}$
○


Regression Planning: Example (2/4)

$$\mathcal{S} = \{a, b, c, d, e\},$$

$$\mathbf{O} = \begin{cases} o_1 = \langle \emptyset, \{b\}, \{-b, c\} \rangle, \\ o_2 = \langle \emptyset, \{e\}, \{b\} \rangle, \\ o_3 = \langle \emptyset, \{c\}, \{b, d, \neg e\} \rangle, \end{cases}$$

$$\mathbf{I} = \{a, b\}$$

$$\mathbf{G} = \{b, d\}$$




Regression Planning: Example (3/4)

$$\mathcal{S} = \{a, b, c, d, e\},$$

$$\mathbf{O} = \begin{cases} o_1 = \langle \emptyset, \{b\}, \{-b, c\} \rangle, \\ o_2 = \langle \emptyset, \{e\}, \{b\} \rangle, \\ o_3 = \langle \emptyset, \{c\}, \{b, d, \neg e\} \rangle, \end{cases}$$

$$\mathbf{I} = \{a, b\}$$

$$\mathbf{G} = \{b, d\}$$




Regression Planning: Example (4/4)

$$\mathcal{S} = \{a, b, c, d, e\},$$

$$\mathbf{O} = \begin{cases} o_1 = \langle \emptyset, \{b\}, \{-b, c\} \rangle, \\ o_2 = \langle \emptyset, \{e\}, \{b\} \rangle, \\ o_3 = \langle \emptyset, \{c\}, \{b, d, \neg e\} \rangle, \end{cases}$$

$$\mathbf{I} = \{a, b\}$$

$$\mathbf{G} = \{b, d\}$$



Other Types of Search

- Of course, other types of search are possible.
- Change perspective: Do not consider the transition system as the space we have to explore, but consider the search through the space of (incomplete) plans:
 - Progression search: Search through the space of plan prefixes
 - Regression search: Search through plan suffixes
- Partial order planning:
 - Search through partially ordered plans by starting with the empty plan and trying to satisfy (sub-)goals by introducing new actions (or using old ones)
 - Make ordering choices only when necessary to resolve conflicts

Computational Complexity

Computational Complexity

The Planning Problem - Formally

Definition (Plan existence problem (PLANEX))

- Instance: $\Pi = \langle S, O, I, G \rangle$.
- Question: Does there exist a plan *triangle* that solves Π , i.e., $Res(I, \Delta) \supseteq G$?

Definition (Bounded plan existence problem (PLANLEN))

- Instance: $\Pi = \langle S, O, I, G \rangle$ and a positive integer n .
- Question: Does there exist a plan Δ of length n or less that solves Π ?

From a practical point of view, also PLANGEN (generating a plan that solves Π) and PLANLENGEN (generating a plan of length n that solves Π) and PLANOPT (generating an optimal plan) are interesting (but at least as hard as the decision problems).

However, easier to deal with decision problems ...

Basic STRIPS with First-Order Terms

The state space for STRIPS with general first-order terms is infinite

Theorem

PLANEX for STRIPS with first-order terms is undecidable.

- Idea of Proof
- We can use function terms to describe (the index of) tape cells of a Turing machine
- We can use operators to describe the Turing machine control
- The existence of a plan is then equivalent to the existence of a successful computation on the Turing machine
- PLANEX for STRIPS with first-order terms can be used to decide the Halting problem

Propositional STRIPS

Theorem

- PLANEX is PSPACE-complete for propositional STRIPS.
- → Membership follows because we can successively guess operators and compute the resulting states (needs only polynomial space)
- → Hardness follows using again a generic reduction from TM acceptance. Instantiate polynomially many tape cells with no possibility to extend the tape (only poly. space, can all be generated in poly. time)
- PLANLEN is also PSPACE-complete (membership is easy, hardness follows by setting $k = 2^{|\Sigma|}$)

Still not good enough: need to restrict our plans even more...

Propositional, Precondition-free STRIPS

Lets go for even easier plans...

Theorem

The problem of deciding plan existence for precondition-free, propositional STRIPS is in P .

Proof

Do a backward greedy plan generation. Choose all operators that make some goals true and that do not make any goals false. Remove the satisfied goals and the operators from further consideration and iterate the step. Continue until all remaining goals are satisfied by the initial state (succeed) or no more operators can be applied (fail).

Theorem

The problem of deciding whether there exists a plan of length k for precondition-free, propositional STRIPS is NP-complete, even if all effects are positive.

Proof

...either trust me .. or trust [Garey and Johnson '79]?

Current Approaches

- In 1992, Kautz and Selman introduced the idea of planning as satisfiability
 - → Encode possible k-step plans as Boolean formulas and use an iterative deepening search approach
- In 1995, Blum and Furst came up with the planning graph approach → iterative deepening approach that prunes the search space using a graph-structure
- In 1996, McDermott proposed to use (again) an heuristic estimator to control the selection of actions, similar to the original GPS idea

Planning Based on Planning Graphs

- Describe possible developments in a graph structure (use only positive effects)
 - Layered graph structure with fact and action levels
 - Fact level (F level): positive atoms (the first level being the initial state)
 - Action level (A level): actions that can be applied using the atoms in the previous fact level
 - Links: precondition and effect links between the two layers
- Record conflicts caused by negative effects and propagate them
- Extract a plan by choosing only non-conflicting parts of the graph (allowing for parallel actions), as long as they don't impact each other
- Parallelism (for non-conflicting actions) is a great boost for the efficiency.

Example Graph (1/5)

$$I = \{at(p1, c), at(p2, s), at(t1, c)\}, G = \{at(p1, s), in(p2, t1)\}$$

$at(p1, c)$

$at(p2, s)$


$at(t1, c)$

 $F0$

Example Graph (2/5)

$$I = \{at(p1, c), at(p2, s), at(t1, c)\}, G = \{at(p1, s), in(p2, t1)\}$$

All applicable actions and their positive effects are included




Example Graph (3/5)

$$I = \{at(p1, c), at(p2, s), at(t1, c)\}, G = \{at(p1, s), in(p2, t1)\}$$

All applicable actions and their positive effects are included

In order to propagate unchanged properties, use noop action, denoted by *




Example Graph (4/5)

$$I = \{at(p1, c), at(p2, s), at(t1, c)\}, G = \{at(p1, s), in(p2, t1)\}$$

All applicable actions and their positive effects are included

In order to propagate unchanged properties, use noop action, denoted by *

Expand graph as long as not all goal atoms are in the fact level




Example Graph (5/5)

$$I = \{at(p1, c), at(p2, s), at(t1, c)\}, G = \{at(p1, s), in(p2, t1)\}$$

All applicable actions and their positive effects are included

In order to propagate unchanged properties, use noop action, denoted by *

Expand graph as long as not all goal atoms are in the fact level

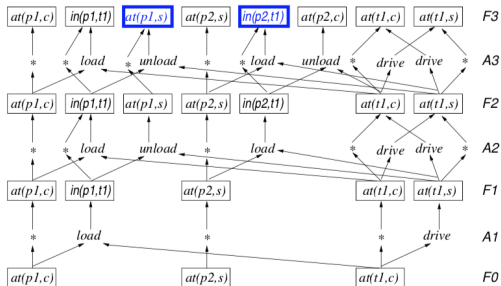


Plan Extraction

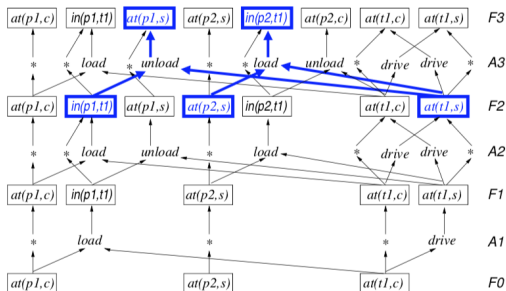
- Start at last fact level with goal atoms
- Select a minimal set of non-conflicting actions that generate the goal atoms
 - Two actions are conflicting if they have complementary effects or if one action deletes or
 - or asserts a precondition of the other action
- Use the preconditions of the selected actions as (sub-)goals on the next lower fact level
- Backtrack if no non-conflicting choice is possible
- If all possibilities are exhausted, the graph has to be extended by another level.

Extracting From the Example Graph (1/5)

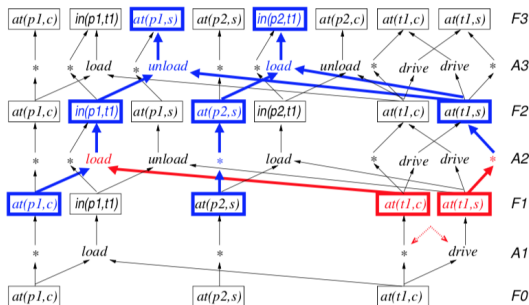
Start with goals at highest fact level




Select minimal set of actions and corresponding subgoals



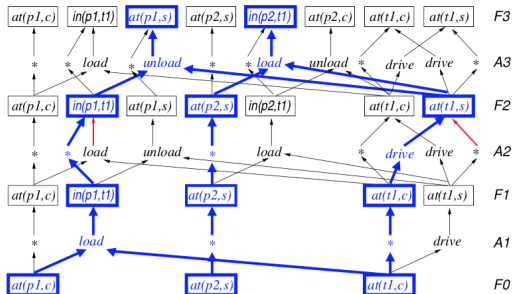
Wrong choice leading to conflicting actions



Other choice, but no further selection possible



Final Selection



Downloaded from <http://ajphaphysocpharm.sagepub.com/> at 11:01 11 November 2014

If a domain contains many symmetries, proving that there is no plan up to length of $k - 1$ can be very costly.


- there is one robot with two grippers
- there is room A that contains n balls
- there is another room B connected to room A the goal is to bring all balls to room B

Obviously, the plan must have a length of at least $n/2$, but ID planners will try out all permutations of actions for shorter plans before noting this.

Prof. Dr. Matthias Schott

Heuristic Search Planning

- Use an heuristic estimator in order to select the next action or state
- Depending on the search scheme and the heuristic, the plan might not be the shortest one
- → It is often easier to go for sub-optimal solutions (remember Logistics)



Heuristic search planner vs. iterative deepening on **Gripper**

- One can use progression or regression search, or even search in the space of incomplete partially ordered plans
- One can use local or global, systematic search strategies
- One can use different heuristics, which can be compared along the dimension of being
 - efficiently computable, i.e., should be computable in poly. time
 - informative, i.e., should make reasonable distinctions between search nodes
 - and admissible, i.e., should underestimate the real costs (useful in A* search).

- Consider all states that are reachable by executing one action
- Try to improve the heuristic value
- Hill climbing: Select the successor with the minimal heuristic value
- Enforced hill climbing: Do a breadth-first search until you find a node that has a better evaluation than the current one.
- → Note: Because these algorithms are not systematic, they cannot be used to prove the absence of a solution

Global Search

2nd type of search : global search

- Maintain a list of open nodes and select always the one which is best according to the heuristic
- Weighted A*: combine estimate $h(S)$ for state S and costs $g(S)$ for reaching S using the weight w with $0 \leq w \leq 1$:

$$f(S) = w * g(S) + (1 - w) * h(S).$$

- If $w = 0.5$, we have ordinary A*, i.e., the algorithm finds the shortest solution provided h is admissible, i.e., the heuristics never overestimates
- If $w < 0.5$, the algorithm is greedy
- If $w > 0.5$, the algorithm behaves more like best-first search

Apart from the search: we need a heuristics...

- General principle for deriving heuristics:
 - Define a simplification (relaxation) of the problem and take the difficulty of a solution for the simplified problem as an heuristic estimator
- Example: straight-line distance on a map to estimate the travel distance
- Example: decomposition of a problem, where the components are solved ignoring the interactions between the components, which may incur additional costs
- In planning, one possibility is to ignore negative effects

Ignoring Negative Effects: Example

In Logistics: The negative effects in load and drive are ignored:

- Simplified load operation: $load(O, V, P)$
 - Precondition: $at(O, P), at(V, P), truck(V)$
 - Effects: $\neg at(\bar{O}, P), in(O, V)$

After loading, the package is still at the place and also inside the truck

- Simplified drive operation: $drive(V, P1, P2)$
 - Precondition: $at(V, P1), truck(V), street(P1, P2)$
 - Effects: $\neg at(\bar{V}, P1), at(V, P2)$

After driving the truck is in two place

→ We want the length of the shortest relaxed plan $h + (s)$

How difficult is monotonic planning?

Monotonic Planning

Assume that all effects are positive

- finding some plan is easy:
 - Iteratively, execute all actions that are executable and have not all their effects made true yet
 - If no action can be executed anymore, check whether the goal is satisfied
 - If not, there is no plan
 - Otherwise, we have a plan containing each action only once
- Finding the shortest plan: easy or difficult?
- PLANLEN for precondition-free operators with only positive effects is NP-complete
- Consider approximations to h^+ .

The HSP Heuristic

- The first idea of estimating the distance to the goal for monotonic planning might be to count the number of unsatisfied goals atoms
 - Neither admissible nor very informative
- Estimate the costs of making an atom p true in state S :

$$h(S, p) = \begin{cases} 0 & \text{if } p \in S \\ \min_{a \in \mathbf{O}, p \in \text{eff}^+(a)} (1 + \max_{q \in \text{pre}(a)} h(S, q)) & \text{otherwise} \end{cases}$$

- Estimate distance from S to S' : $h(S, S') = \max_{p \in S'} h(S, p)$
- Is admissible, because only the longest chain is taken, but it is not very informative
- Use "sum" instead of "max" (this is the HSP heuristics)
- Is not admissible, but more informative. However, it ignores positive interactions!
- \rightarrow Can be computed by using a dynamic programming technique

Summary

- Rational agents need to plan their course of action
- In order to describe planning tasks in a domain-independent, declarative way, one needs planning formalisms
- Basic STRIPS is a simple planning formalism, where actions are described by their preconditions in form of a conjunction of atoms and the effects are described by a list of literals that become true and false
- PDDL is the current "standard language" that has been developed in connection with the international planning competition
- Basic planning algorithms search through the space created by the transition system or through the plan space.
- Planning with STRIPS using first-order terms is undecidable
- Planning with propositional STRIPS is PSPACE-complete
- Since 1992, we have reasonably efficient planning method for propositional, classical STRIPS planning